Asynchronous Message Delivery System

Biagio Cornacchia
Matteo Abaterusso

Overview

- A server that **subscribes** to the system obtains a **virtual IP** and **virtual MAC**
- When the server is online, all client requests to that virtual IP are delivered **synchronously**
- When the server is offline, the system **stores** all requests destinated to it
- When the server **comes back online**, packets stored inside the controller are immediately sent out to it
- A server can subscribe/unsubscribe or change its status using the **Restful APIs** exposed by the system

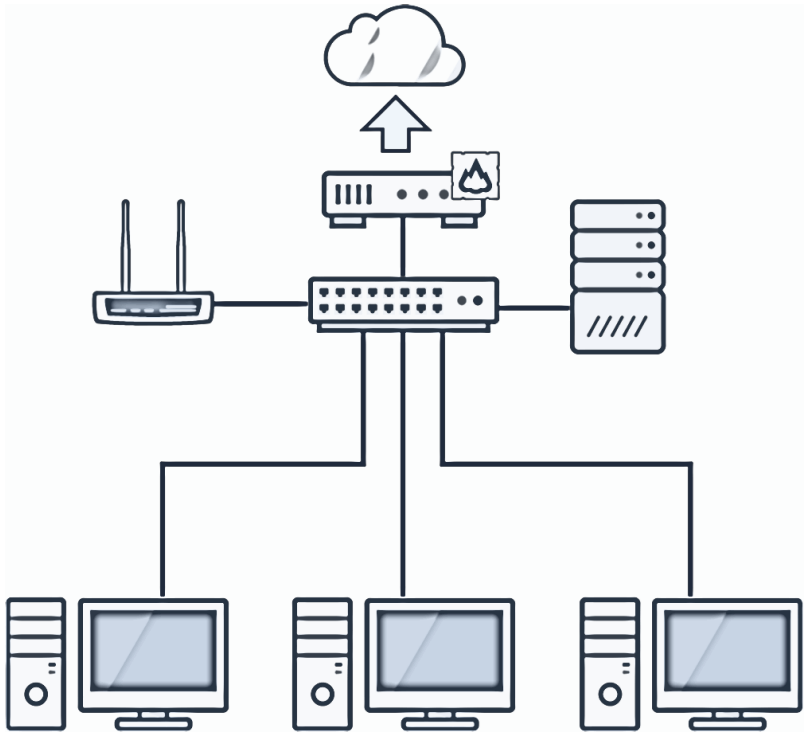

Design

Restful APIs

Subscription

http://<CONTROLLER_IP>:8080/amd/server/subscription/json

- POST
- DELETE

Status

http://<CONTROLLER_IP>:8080/amd/server/status/json

- PUT

Info

http://<CONTROLLER_IP>:8080/amd/server/info/json

- GET

ARP Requests

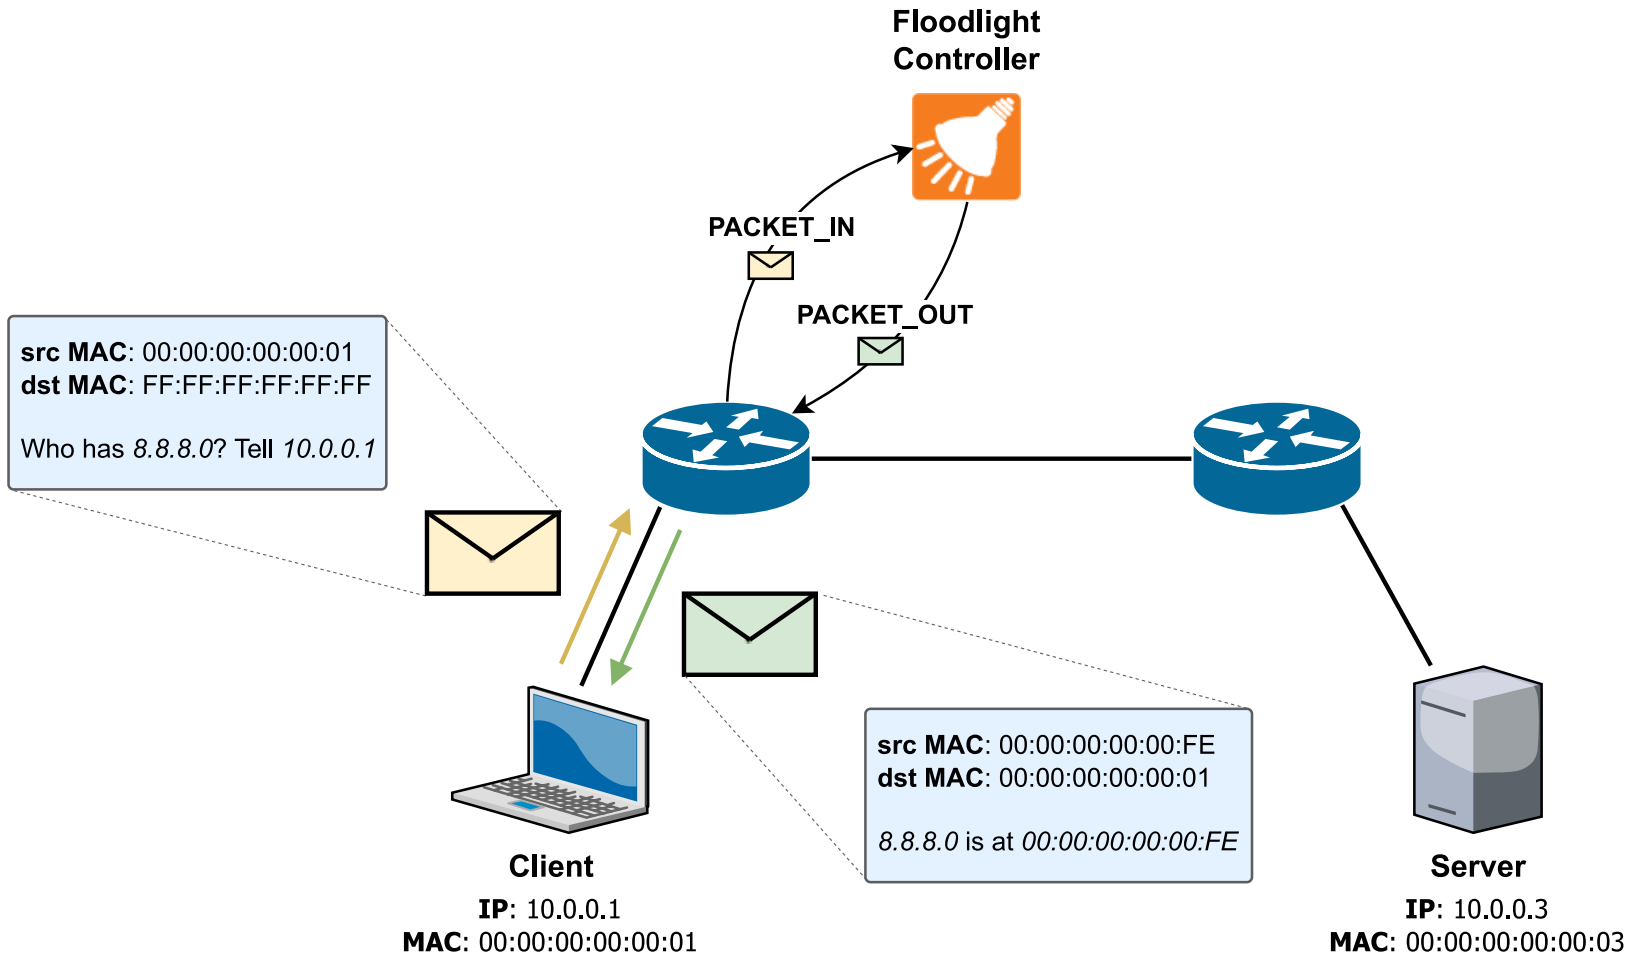

IPv4 Requests (Server online)

- **Client** to **Server** path

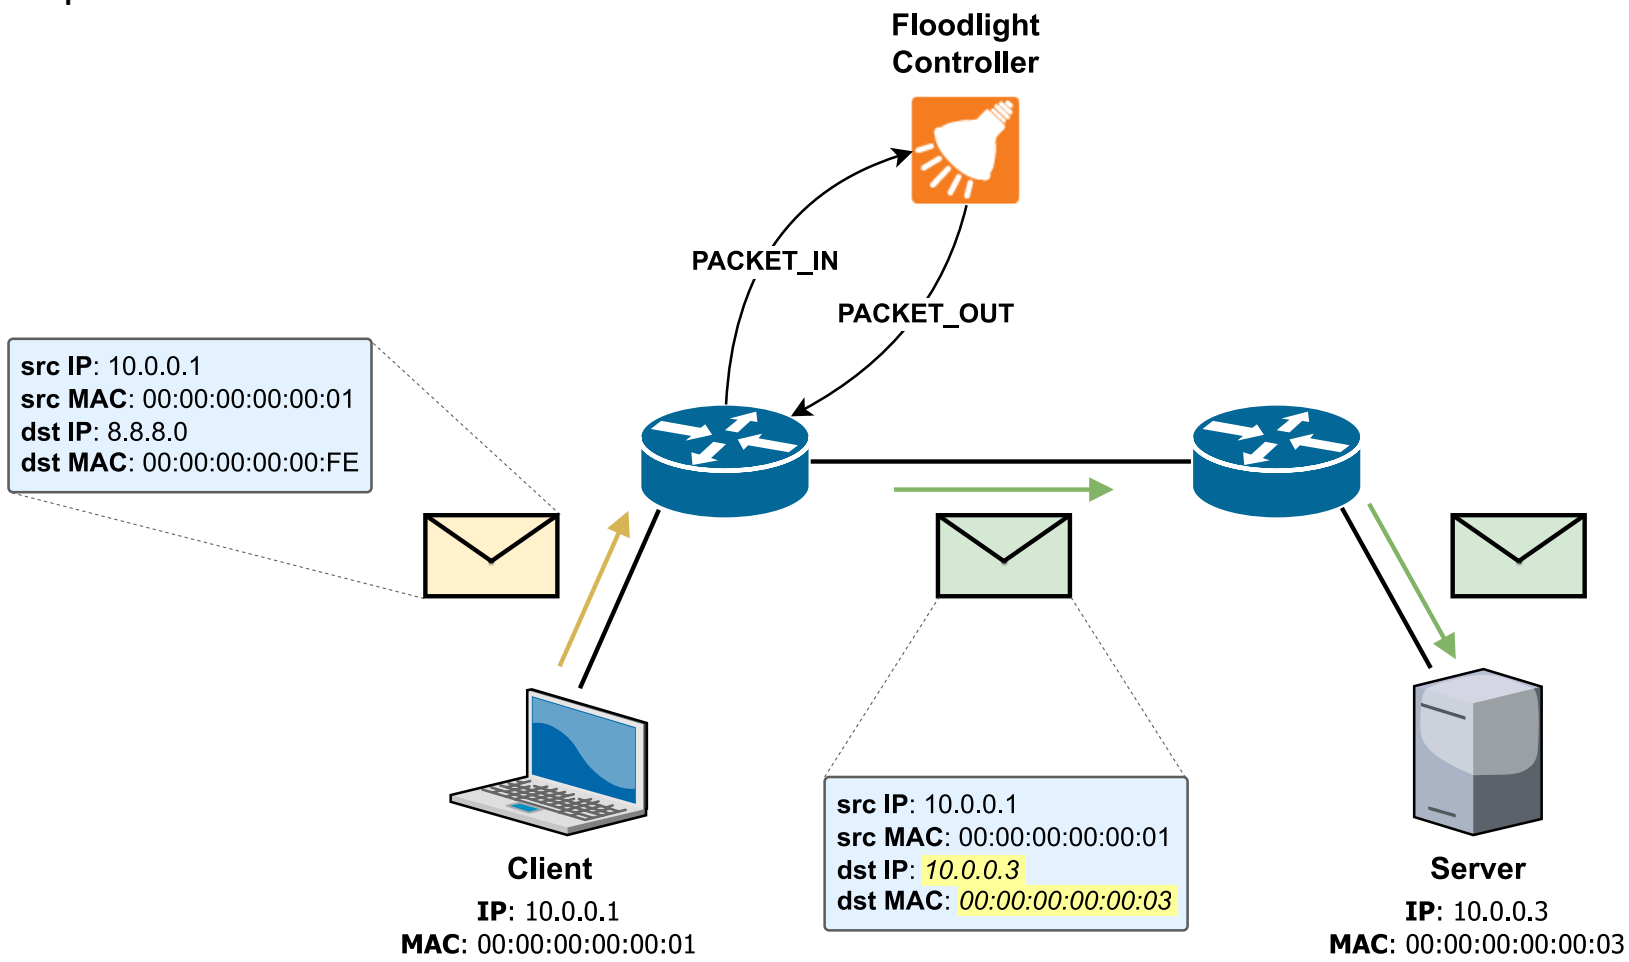

IPv4 Requests (Server online)

- **Server** to **Client** path

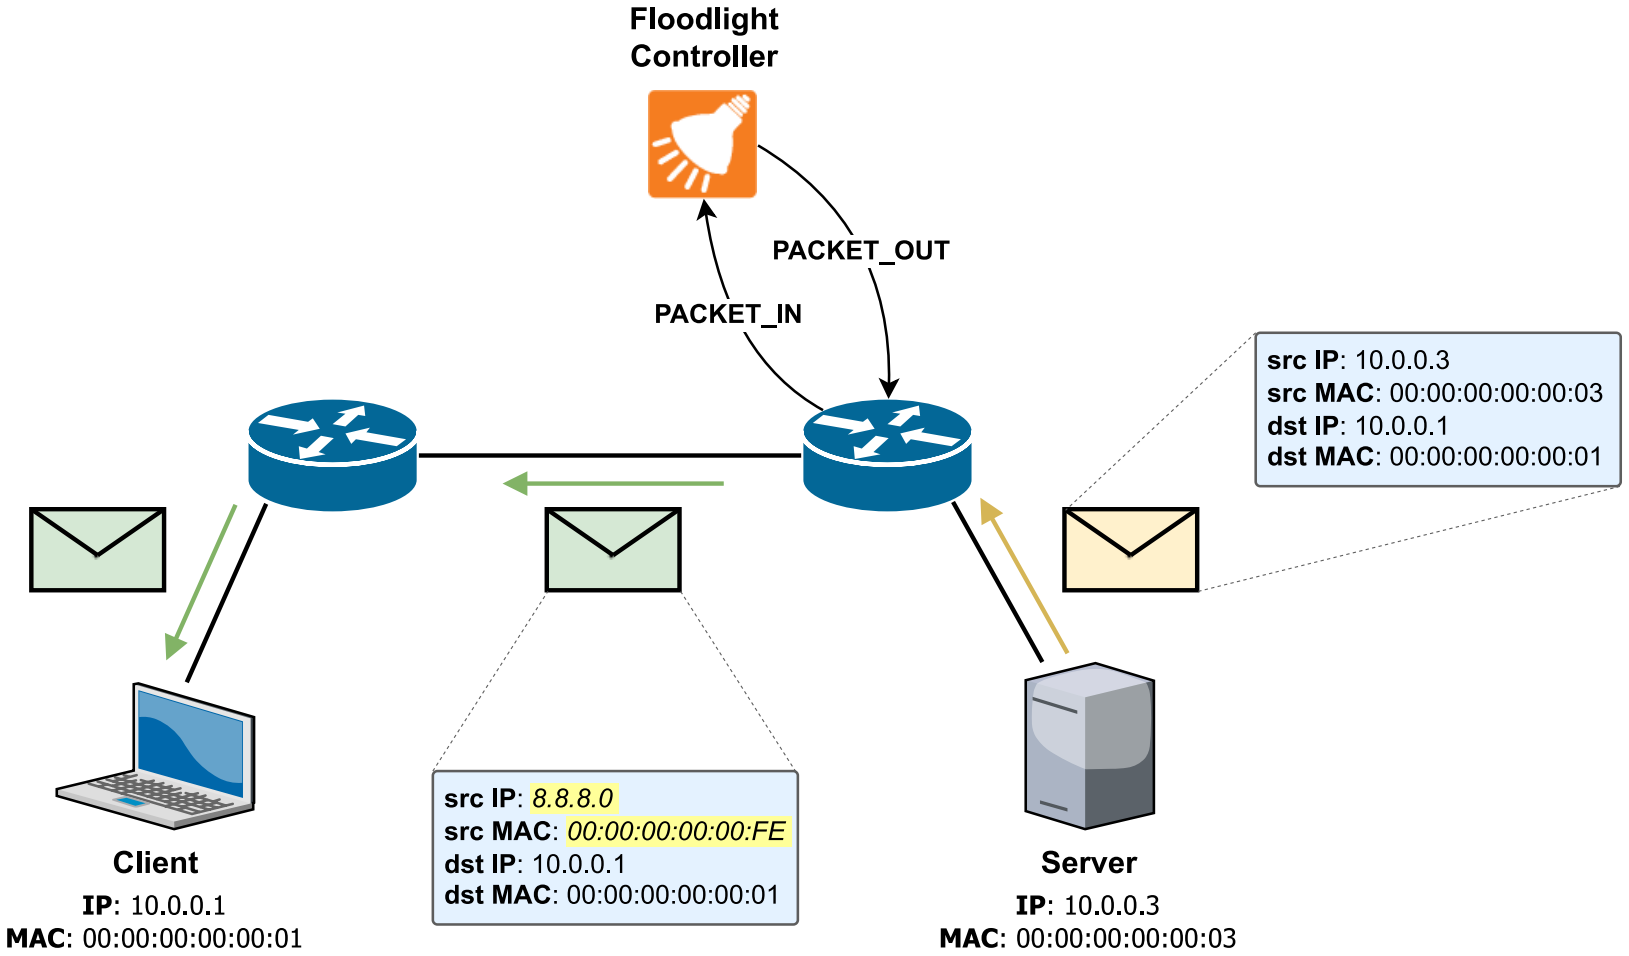

IPv4 Requests (Server offline)

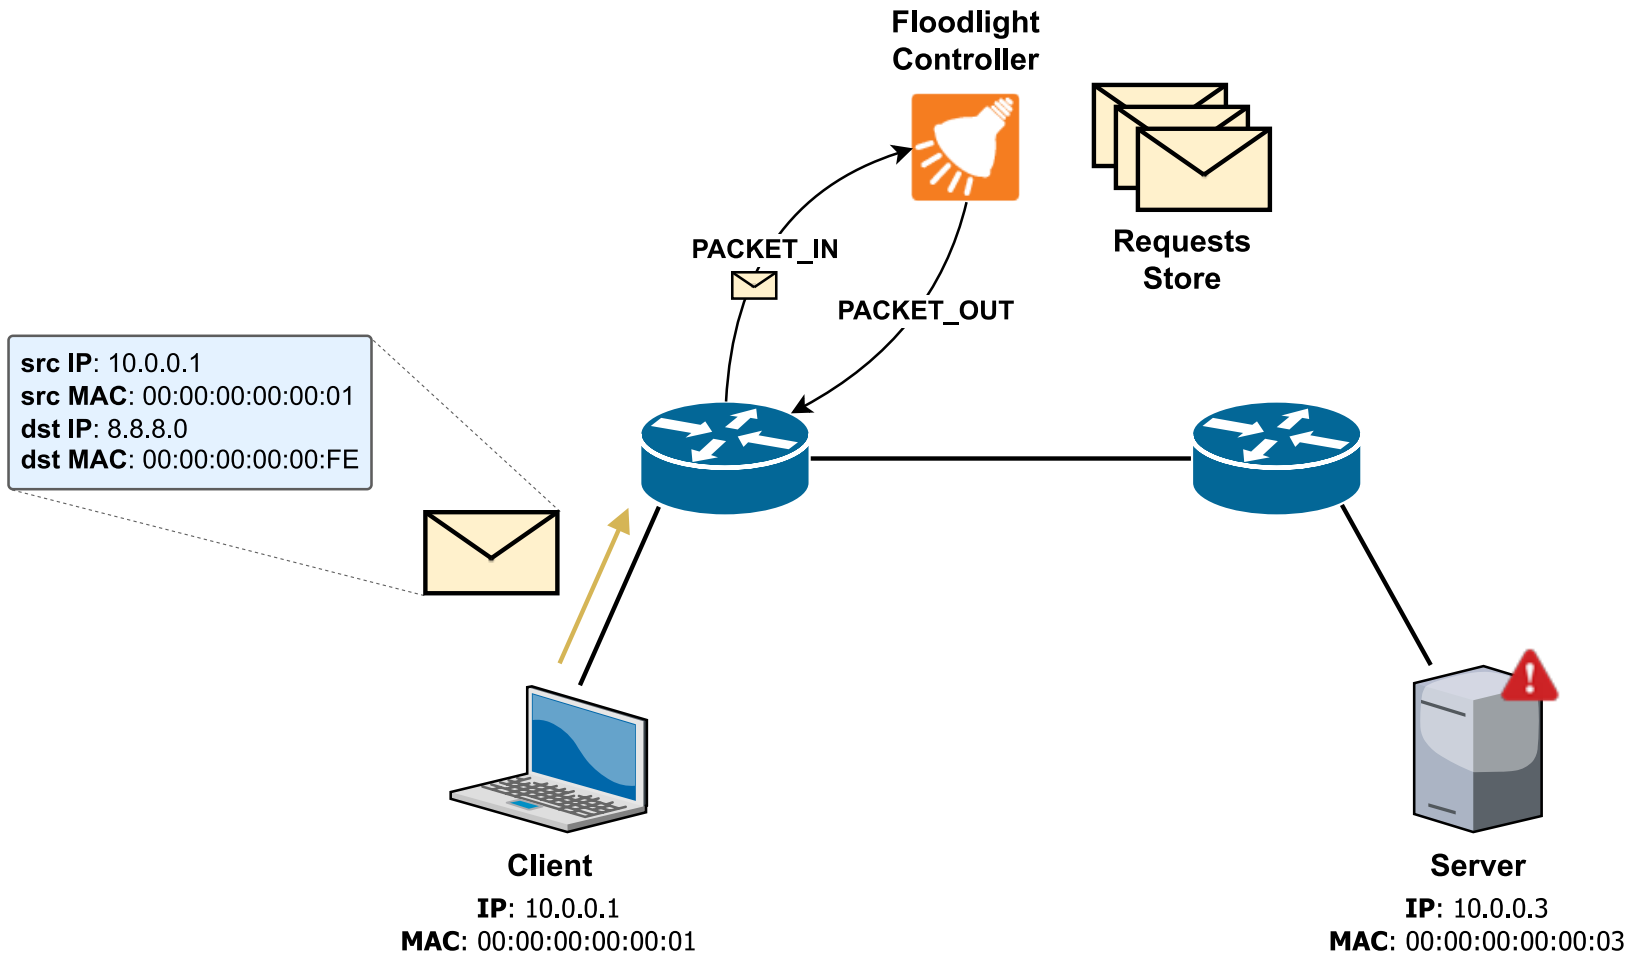

Flow Entries

- The **soft timeout** and **hard timeout** establish the lifetime of the flow entries on a switch.
- This will cause a **transition time** during which packets will be lost. So, there is a **trade-off** on the timeout choice.
- The chosen approach is to **actively intervene** when the server status changes by removing the flow entries related to that server.
- The fact that translation flow entries are installed only on **access switches** is exploited to interact with a reasonable number of switches.

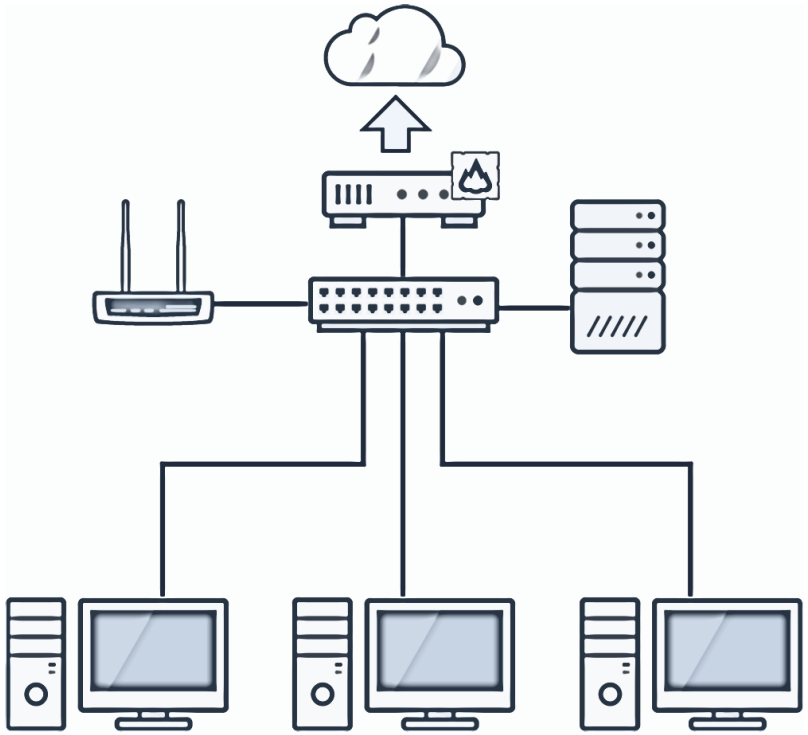

Implementation

Forwarding Module Problem

- In the Floodlight pipeline the **Forwarding** module is located before the **AMD** module.
- The **default behavior** of the Forwarding module is that if the destination address of the packet is **unknown** then a **flooding** action is applied, otherwise the packet is **directly forwarded**.
- This results in having **duplicated packets** in the network.

Forwarding Module Solution

- In the AMD module

```
@Override
public boolean isHandledByAMD(String srcIP, String dstIP) {
    // Check if a packet is related to a subscribed server and should be translated
    for (String[] addresses : virtualAddresses) {
        if (dstIP.equals(addresses[1]) || (subscribedServers.containsKey(addresses[1])
                && srcIP.equals(subscribedServers.get(addresses[1]).getIPAddress()))) {
            logger.info("(FORWARDER) IPv4 packet managed by AMD");
            return true;
        }
    }

    return false;
}
```

Forwarding Module Solution

- In the processPacketInMessage method of the Forwarding class

```
@Override
public Command processPacketInMessage(IOFSwitch sw, OFPacketIn pi, IRoutingDecision decision, FloodlightContext cntx) {
    Ethernet eth = IFloodlightProviderService.bcStore.get(cntx, IFloodlightProviderService.CONTEXT_PI_PAYLOAD);

    OFPort inPort = OFMessageUtils.getInPort(pi);
    NodePortTuple npt = new NodePortTuple(sw.getId(), inPort);

    if (eth.getPayload() instanceof IPv4) {
        IPv4 pkt = (IPv4)eth.getPayload();

        IAsynchronousMessageDelivery amd = this.context.getServiceImpl(IAsynchronousMessageDelivery.class);
        if (amd.isHandledByAMD(pkt.getSourceAddress().toString(), pkt.getDestinationAddress().toString())) {
            return Command.CONTINUE;
        }
    }
```

OVS Security Policy Problem

By default, the **Open vSwitch** implements a security policy that **drops** packets with an **external source MAC address**.

The parameters specified in the flow mod are the following:

- **priority** = 1
- **dl src** = VIRTUAL MAC
- **action** = normal

Flow Mod Optimization

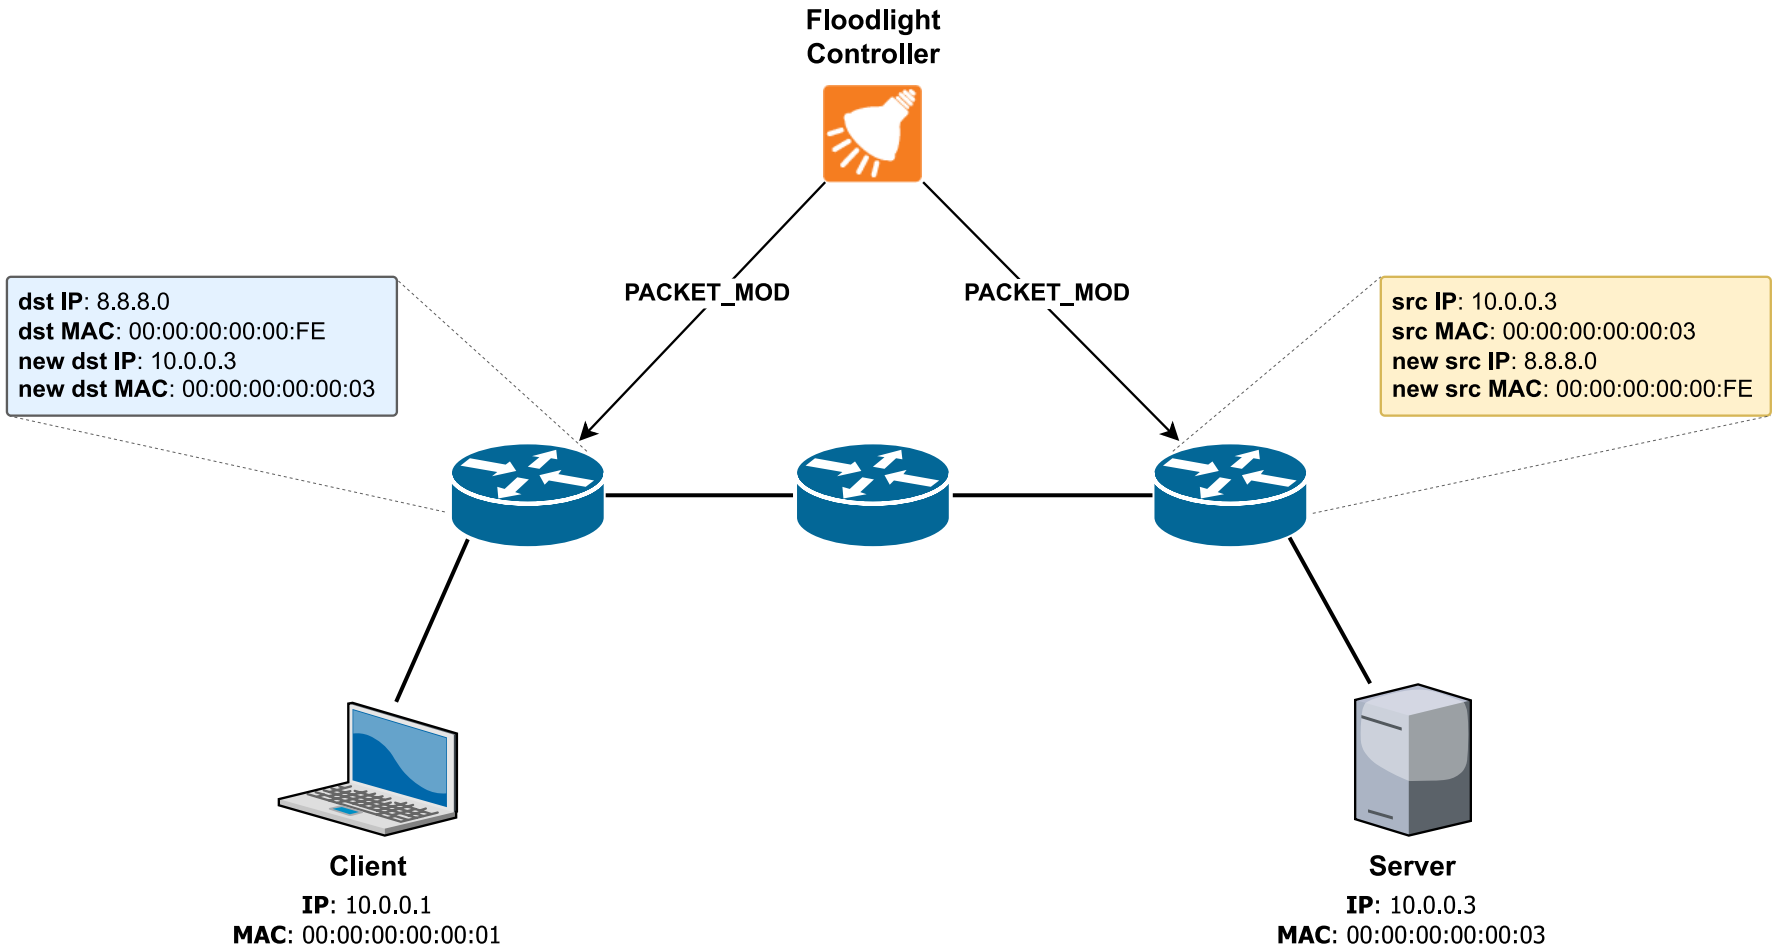

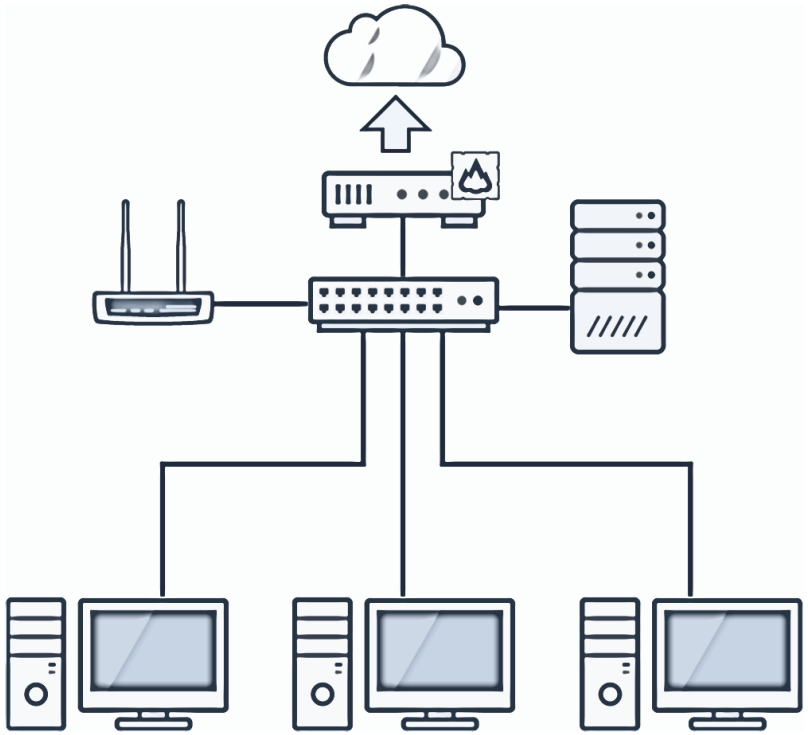

Testing

Testing Network

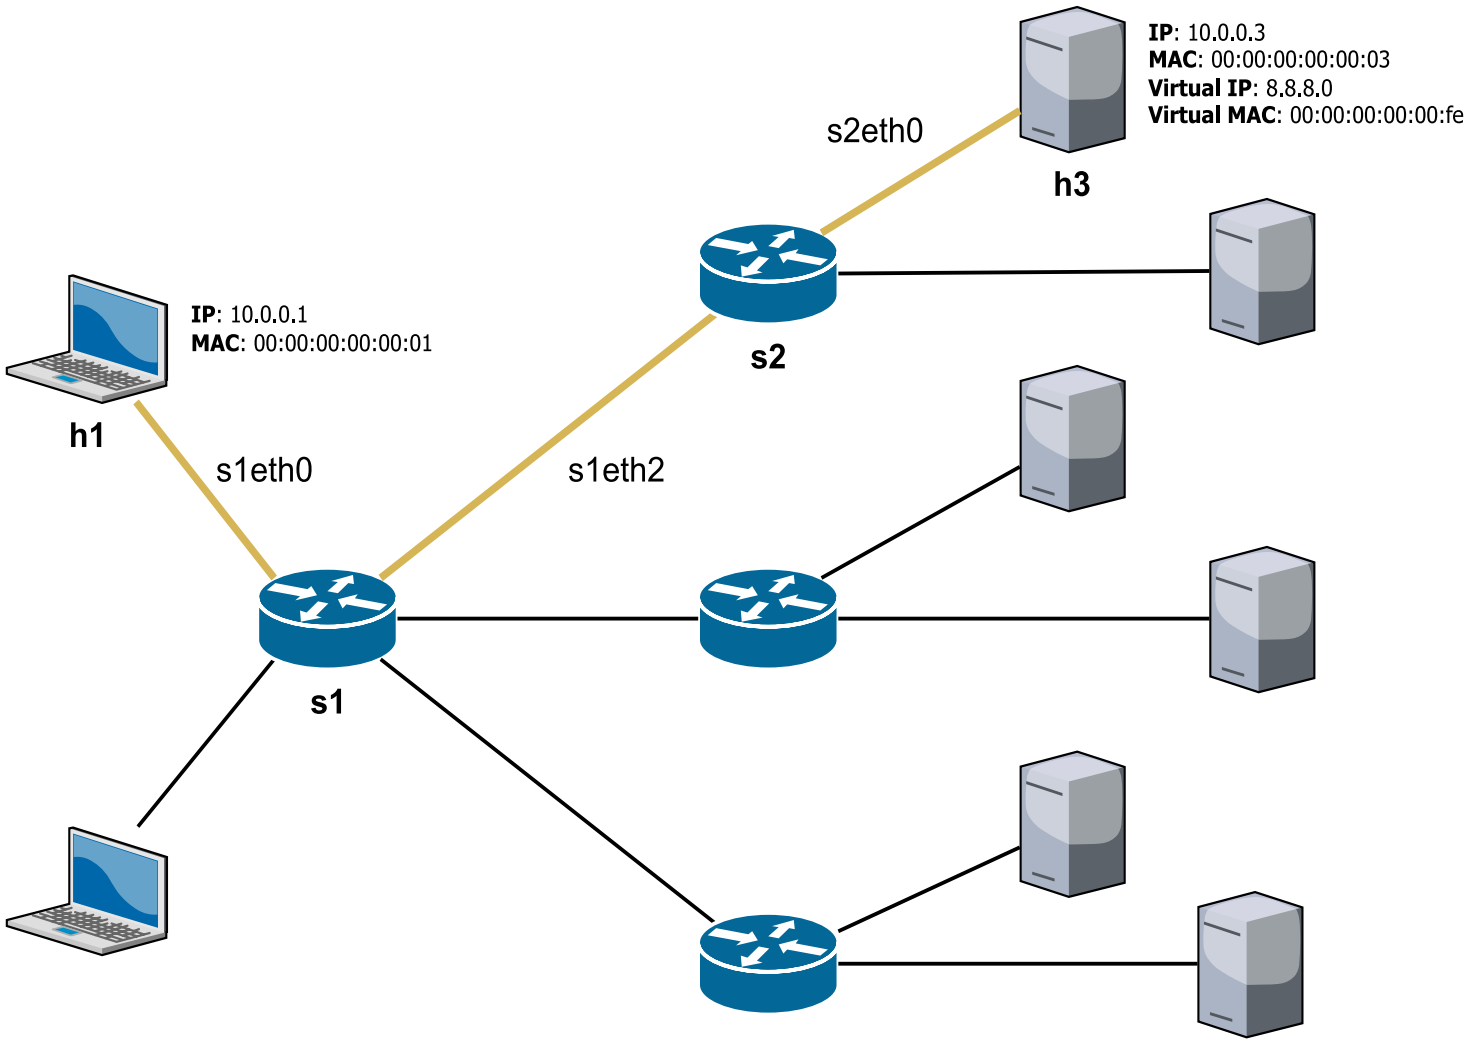

Server Subscription

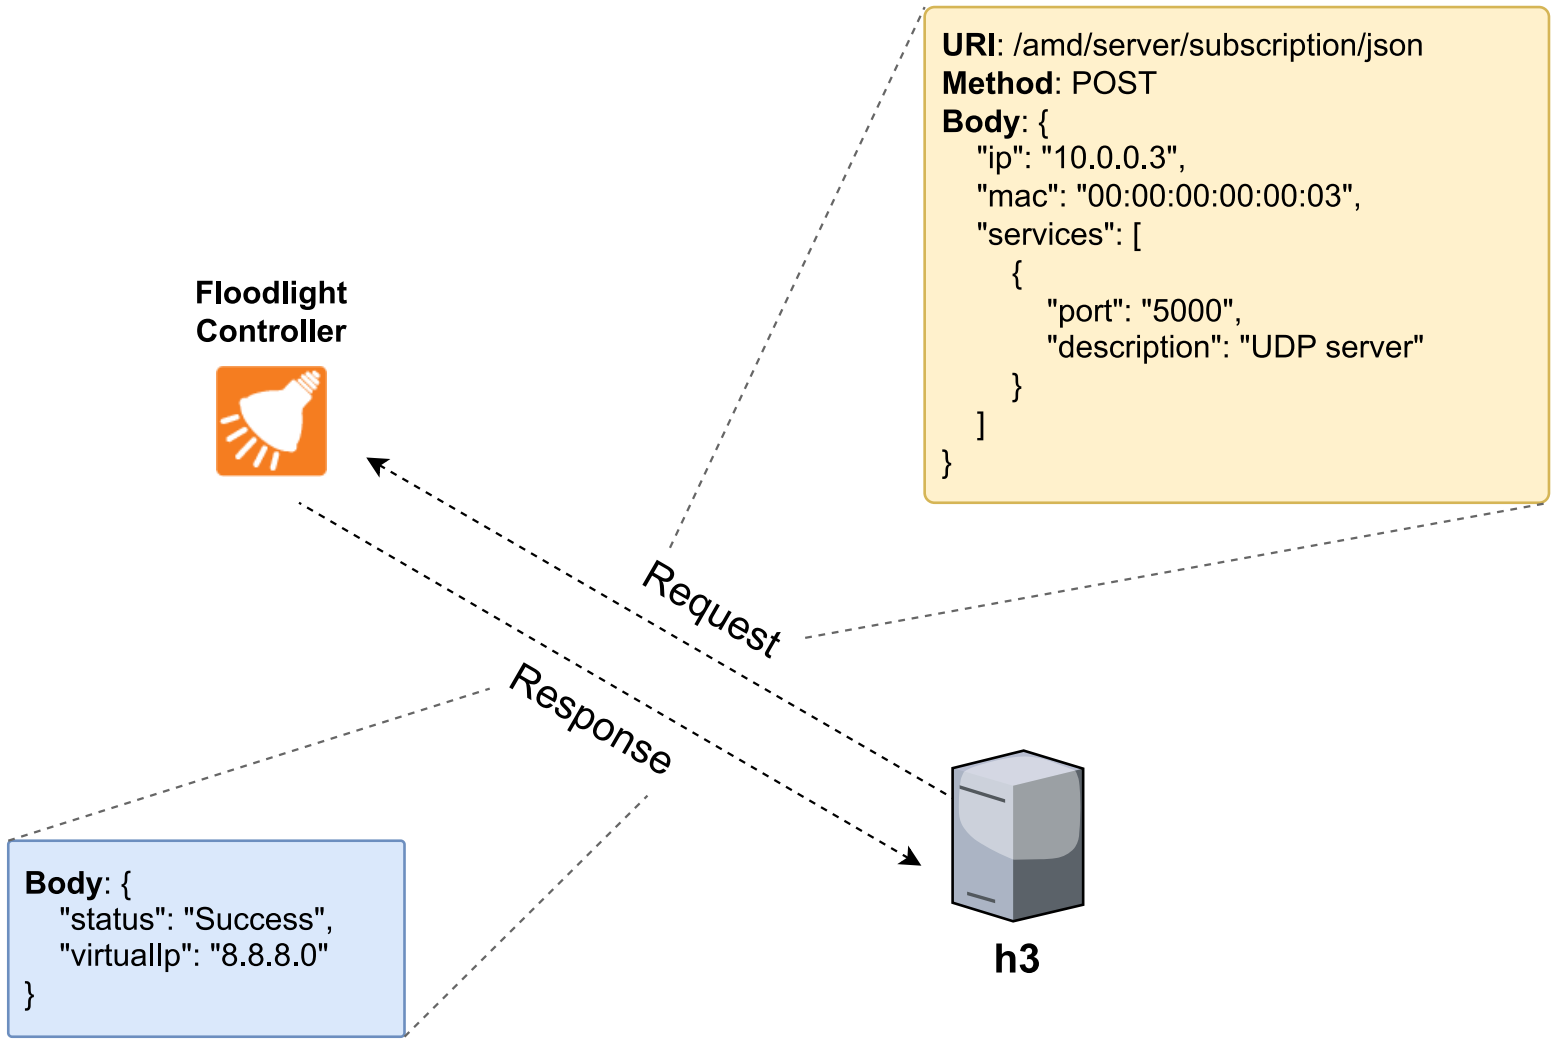

Server Info

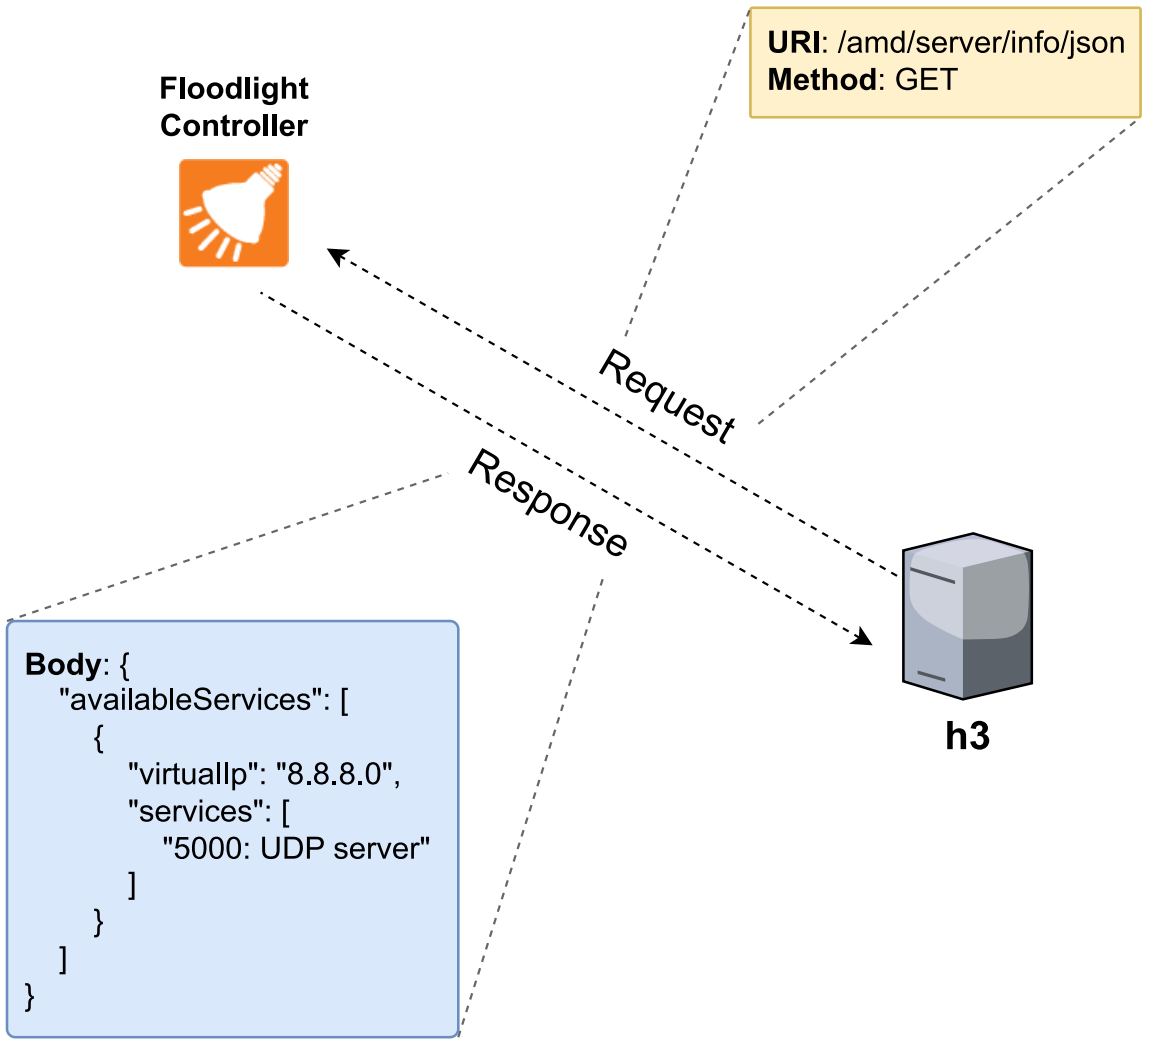
Floodlight
Controller
URI: /amd/server/info/json
Method: GET
Request
Response
h3
Body: {
"availableServices": [
{
"virtualIp": "8.8.8.0",
"services": [
"5000: UDP server"
]
}
]
}

Server Status Update

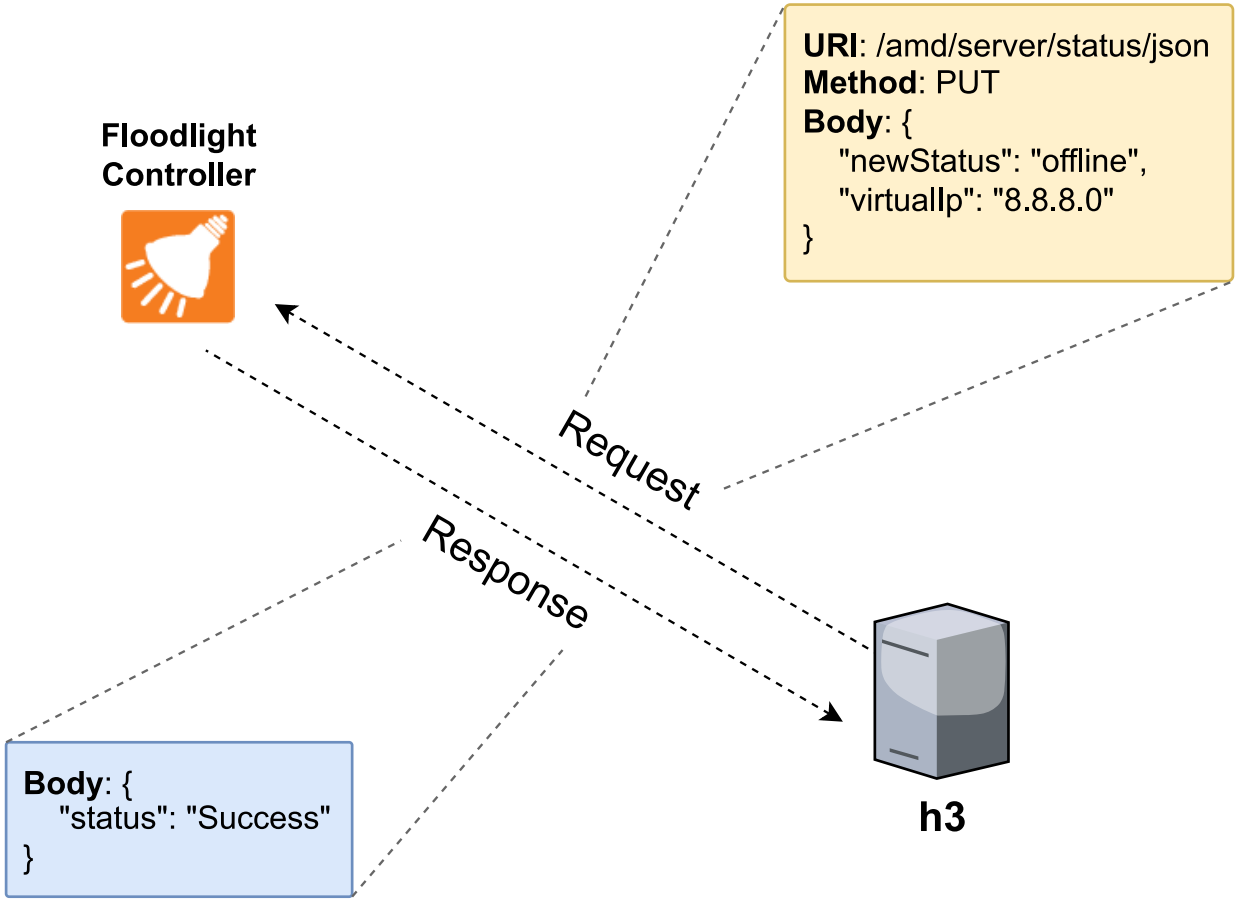

Server Status Update

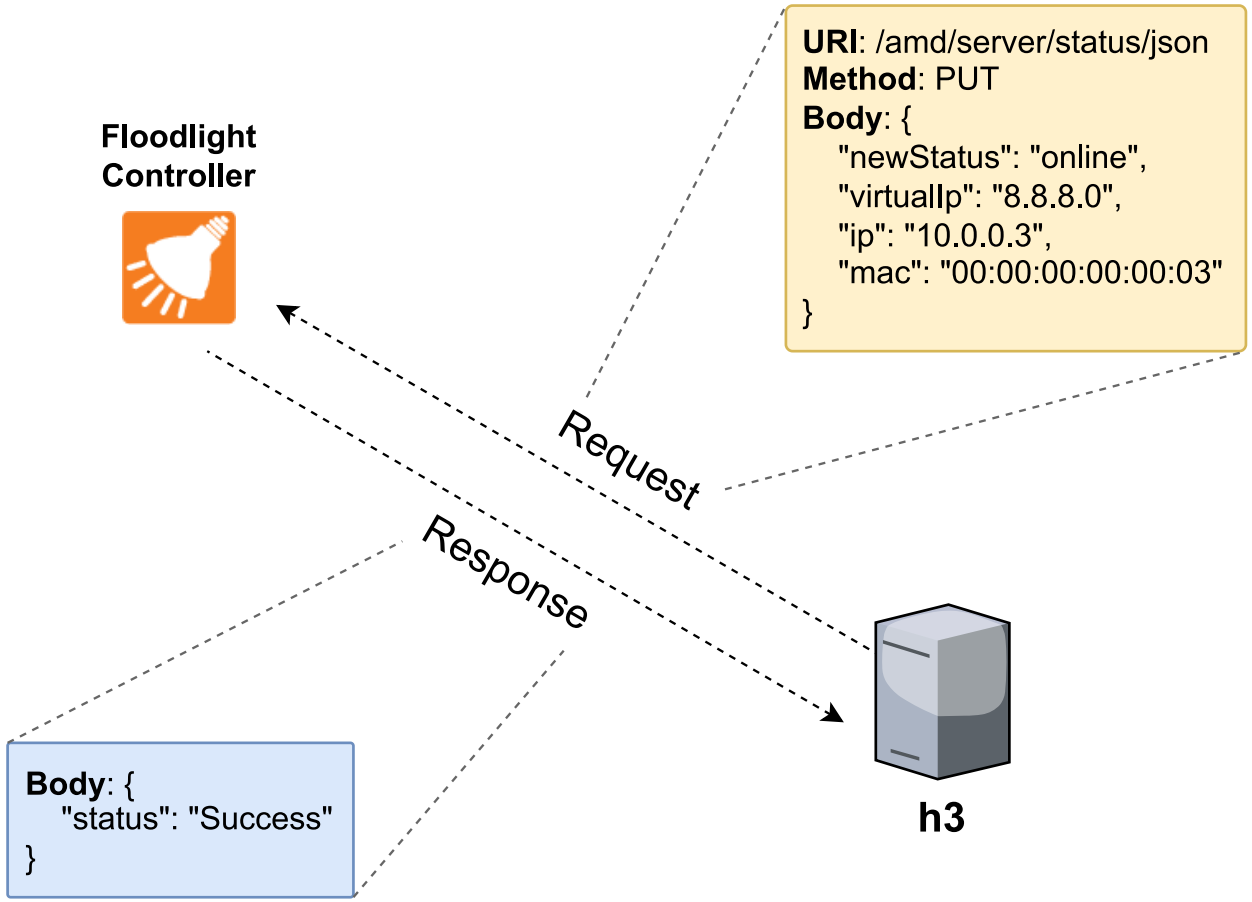

Server Unsubscription